

# Linear classifiers. Kernels

V. Kitov



Yandex School of Data Analysis

Imperial College London  
Department of Physics

January 2015

# Table of Contents

- 1 Stochastic gradient descent
- 2 Regularization
- 3 Logistic regression
- 4 Support vector machines
- 5 Kernel support vector machines

## Linear discriminant functions

- Linear discriminant function:  $g(x) = w^T x + w_0$ ,

$$\hat{\omega} = \begin{cases} \omega_1, & g(x) \geq 0 \\ \omega_2, & g(x) < 0 \end{cases}$$

- If we denote classes  $\omega_1$  and  $\omega_2$  with  $y = +1$  and  $y = -1$  respectively, we get the decision rule  $y = \text{sign } g(x)$ .
- Define new feature  $x_0 \equiv 1$ , then  $g(x) = w^T x = \langle w, x \rangle$  for  $w = [w_0, w_1, \dots, w_D]^T$ .
- Define margin  $M(x) = g(x)y$ 
  - $M(x) \geq 0 \iff$  object  $x$  is correctly classified
  - $|M(x)|$  measures confidence of decision

## Weights selection problem

- Final task - minimize misclassifications count:

$$Q_{\text{accurate}}(w|X) = \sum_i \mathbb{I}[M(x_i|w) < 0] \rightarrow \min_w$$


- Standard optimization techniques are impossible, because  $Q(w, X)$  is discontinuous.
- Idea: approximate indicator of misclassification with smooth majorizing function  $\mathcal{L}$ :

$$\mathbb{I}[M(x_i|w) < 0] \leq \mathcal{L}(M(x_i|w))$$

## Approximation of target criteria

We obtain approximation of target criteria:

$$\begin{aligned} Q_{\text{accurate}}(w|X) &= \sum_i \mathbb{I}[M(x_i|w) < 0] \\ &\leq \sum_i \mathcal{L}(M(x_i|w)) = Q_{\text{approx}}(w|X) \end{aligned}$$



$$\begin{aligned} Q(M) &= (1 - M)^2 \\ V(M) &= (1 - M)_+^{-1} \\ S(M) &= 2(1 + e^M)^{-1} \\ L(M) &= \log_2(1 + e^{-M}) \\ E(M) &= e^{-M} \end{aligned}$$

# Optimization

- Optimization task to get weights:

$$Q_{approx}(w|X) = \sum_{i=1}^n \mathcal{L}(M(x_i|w)) = \sum_{i=1}^n \mathcal{L}(\langle w, x_i \rangle y_i) \rightarrow \min_w$$

- Gradient descent algorithm:
  - Iteratively until convergence

$$w \leftarrow w - \eta \frac{\partial Q_{approx}(w|X)}{\partial w} = w - \eta \sum_{i=1}^n \mathcal{L}'(\langle w, x_i \rangle y_i) x_i y_i$$

- $\eta$  - parameter, controlling the speed of convergence.
- Faster convergence when updates are more often - e.g. at each observation. Observations may be taken randomly.

## Improved optimization

### Stochastic gradient descent algorithm

Calculate  $\hat{Q}_{approx}(w, X) = \sum_{i=1}^n \mathcal{L}(M(x_i|w))$

Iteratively, until convergence of  $\hat{Q}_{approx}$  or convergence of  $w$ :

- 1 select random observation  $(x_i, y_i)$
- 2 adapt weights:  $w \leftarrow w - \eta \mathcal{L}'(\langle w, x_i \rangle y_i) x_i y_i$
- 3 Estimate error:  $\varepsilon_i = \mathcal{L}(\langle w, x_i \rangle y_i)$
- 4 Recalculate  $\hat{Q}_{approx} = (1 - \alpha) \hat{Q}_{approx} + \alpha \varepsilon_i$

Initial weights selection:

- all zeros
- random at  $[-\frac{1}{2D}, \frac{1}{2D}]$  (for logistic approximation) or arbitrary random
- $w_i = \frac{\langle x^i, y \rangle}{\langle x^i, x^i \rangle}$

# Analysis

## Advantages

- Easy to implement
- Works in online environments
- Small random subset of objects may be enough for accurate complete estimation



# Analysis

## Advantages

- Easy to implement
- Works in online environments
- Small random subset of objects may be enough for accurate complete estimation

## Disadvantages

- May converge to local optima
- For improper choice of parameters
  - may diverge
  - may converge too slowly
- for large  $D$  and small  $n$  may overtrain
- when  $\mathcal{L}(u)$  has horizontal asymptotes, algorithm may get stuck for large values of  $\langle w, x_i \rangle$

## Examples

Delta-rule  $\mathcal{L}(u) = (u - 1)^2$

$$w \leftarrow w - \eta(\langle w, x_i \rangle - y_i)x_i$$

This also fits for regression  $y \in \mathbb{R}$ ,  $a(x) = \langle w, x \rangle$  and cost function  $(\langle w, x \rangle - y)^2$

Perceptron of Rosenblatt  $\mathcal{L}(u) = [-u]_+$

$$w \leftarrow w + \begin{cases} 0, & \langle w, x_i \rangle y_i \geq 0 \\ \eta x_i y_i & \langle w, x_i \rangle y_i < 0 \end{cases}$$

## Recommendations for usage

- Faster converges for scaled features
  - normalization equalizes long narrow valley structures
  - $\langle w, x_i \rangle y_i$  becomes limited at early iterations - SGD does not “get stuck” for  $\mathcal{L}$  with horizontal asymptotes.
- Faster convergence when more errors are made:
  - random sampling with probabilities proportional to  $\varepsilon_i = \mathcal{L}(\langle w, x_i \rangle y_i)$
  - random sampling with most diverse objects (e.g. sampling repeatedly from different classes)
- Faster calculation: make change to  $w$  only for mistakes large enough if  $\varepsilon_i \geq \delta$ , for some threshold  $\delta > 0$ .
- Find global minimum by starting the procedure from different starting points

## Selection of $\eta$

- Larger  $\eta \Rightarrow$  algorithm more prone to diverge.
- Plot  $Q_{approx}(w)$  (or  $\hat{Q}_{approx}(w)$ ) versus iteration number  $t$  to control convergence.
- Deterministic scheme:
  - Stochastic gradient descent converges to local optima if
    - $\eta_t \rightarrow 0$
    - $\sum_{t=1}^{\infty} \eta_t = \infty$
    - $\sum_{t=1}^{\infty} \eta_t^2 < \infty$
  - Example:  $\eta_t = \frac{1}{t}$
- Data dependent scheme:
  - At each step find  $\eta_t = \arg \min_{\eta} Q_{approx}(w - \eta \frac{\partial Q_{approx}}{\partial w})$
  - Often analytical solution for such  $\eta$  exists

# Overtraining

- Early stopping:
  - control algorithm performance on separate validation set.
  - when performance start to increase - stop.
- Regularization
  - Add penalty for large weights:

$$Q_{approx}^{regularized}(w) = Q_{approx}(w) + \frac{\tau}{2}|w|^2$$

- Gradient descent step becomes:  $w \leftarrow w(1 - \eta\tau) - \eta Q'_{approx}(w)$
- Weights get exponential decay at each step
- $\tau$  controls the trade-off between bias and variance
  - it prevents overfitting,
  - prevents non-stable estimates of  $w$  for correlated features
  - prevents SGD getting stuck for large weights and small  $\mathcal{L}'(u)$
  - limits flexibility of the model

# Table of Contents

- 1 Stochastic gradient descent
- 2 Regularization**
- 3 Logistic regression
- 4 Support vector machines
- 5 Kernel support vector machines

# Regularization

- Useful technique to control the trade-off between bias and variance, can be applied to any algorithm.

$$Q^{regularized}(w) = Q(w) + \tau \|w\|_2$$

$$Q^{regularized}(w) = Q(w) + \tau \|w\|_1$$

$$\|w\|_1 = \sum_{d=1}^D |w^d|, \quad \|w\|_2 = \sqrt{\sum_{d=1}^D (w^d)^2}$$

- Examples:
  - LASSO: least-squares regression, using  $\|w\|_1$
  - Ridge: least-squares regression, using  $\|w\|_2$
  - Elastic Net: : least-squares regression, using both

## Maximum probability estimation

- $X = \{x_1, x_2, \dots, x_n\}$ ,  $Y = \{y_1, y_2, \dots, y_n\}$  - training sample of i.i.d. observations,  $(x_i, y_i) \sim p(y|x, w)$
- ML estimation  $\hat{w} = \arg \max_w p(Y|X, w)$
- Using independence assumption:

$$\prod_{i=1}^n p(y_i|x_i, w) = \sum_{i=1}^n \ln p(y_i|x_i, w) \rightarrow \max_w$$

- Approximated misclassification:

$$\sum_{i=1}^n \mathcal{L}(g(x_i)y_i|w) \rightarrow \min_w$$

- Interrelation:

$$\mathcal{L}(g(x_i)y_i|w) = -\ln p(y_i|x_i, w)$$



## Maximum a posteriori estimation

- $X = \{x_1, x_2, \dots, x_n\}$ ,  $Y = \{y_1, y_2, \dots, y_n\}$  - training sample of i.i.d. observations,  $(x_i, y_i) \sim p(x, y|w)$
- $x_i \sim p(x|w)$
- MAP estimation:
  - $w$  is random with prior probability  $p(w)$

$$p(w|X, Y) = \frac{p(X, Y, w)}{p(X, Y)} = \frac{p(X, Y|w)p(w)}{p(X, Y)} \propto p(X, Y|w)p(w)$$

$$w = \arg \max_w p(w|X, Y) = \arg \max_w p(X, Y|w)p(w)$$

$$\sum_{i=1}^n \ln p(x_i, y_i|\theta) + \ln p(w) \rightarrow \max_w$$

## Gaussian prior

- Gaussian prior

$$\ln p(w, \sigma^2) = \ln \left( \frac{1}{(2\pi\sigma^2)^{n/2}} e^{-\frac{\|w\|_2^2}{2\sigma^2}} \right) = -\frac{1}{2\sigma^2} \|w\|_2^2 + \text{const}(w)$$

- Laplace prior

$$\ln p(w, C) = \ln \left( \frac{1}{(2C)^n} e^{-\frac{\|w\|_1}{C}} \right) = -\frac{1}{C} \|w\|_1 + \text{const}(w)$$

$L_1$  norm

- $\|w\|_1$  regularizer will do feature selection.
- Consider

$$Q(w) = \sum_{i=1}^n \mathcal{L}_i(w) + \frac{1}{C} \sum_{d=1}^D |w_d|$$

- if  $\frac{1}{C} > \sup_w \left| \frac{\partial \mathcal{L}(w)}{\partial w_i} \right|$ , then it becomes optimal to set  $w_i = 0$
- For smaller  $C$  more inequalities will become active.

## Adaptive feature importances

- Suppose
  - weights have prior Gaussian distribution
  - are uncorrelated
  - each weight  $w_d$  has individual prior variance  $\sigma_d^2 = C_d$
- Prior distribution becomes:

$$p(w) = \frac{1}{(2\pi)^{n/2} \sqrt{C_1 \dots C_D}} e^{-\sum_{d=1}^D \frac{w_d^2}{2C_d}}$$

- Target functional becomes:

$$Q_{approx}(w) = \sum_{i=1}^n \mathcal{L}_i(w) + \frac{1}{2} \sum_{d=1}^D \left( \ln C_d + \frac{w_d^2}{C_d} \right) \rightarrow \min_{w, C}$$

- If  $\hat{C}_d \rightarrow 0$ , feature  $d$  is removed.

# Table of Contents

- 1 Stochastic gradient descent
- 2 Regularization
- 3 Logistic regression**
- 4 Support vector machines
- 5 Kernel support vector machines

# Logistic regression

Assume ( $\gamma_1, \gamma_2$  are the costs of misclassifying classes  $\omega_1$  and  $\omega_2$ ):

$$\ln \left( \frac{\gamma_1 p(\omega_1|x)}{\gamma_2 p(\omega_2|x)} \right) = \beta_0 + \beta^T \mathbf{x}$$

It is equivalent to

$$\begin{aligned} p(\omega_2|x) &= \frac{1}{1 + \exp(\beta'_0 + \beta^T \mathbf{x})} \\ p(\omega_1|x) &= \frac{\exp(\beta'_0 + \beta^T \mathbf{x})}{1 + \exp(\beta'_0 + \beta^T \mathbf{x})} \end{aligned}$$

where  $\beta'_0 = \beta_0 - \ln(\gamma_1/\gamma_2)$

# Logistic regression

Decision rule (following Bayes minimum risk principle):

$$\mathbf{x} = \begin{cases} \omega_1, & \beta'_0 + \boldsymbol{\beta}^T \mathbf{x} > 0 \\ \omega_2, & \beta'_0 + \boldsymbol{\beta}^T \mathbf{x} < 0 \end{cases}$$

Estimate with ML:

$$\prod_{i=1}^n p(c_i | x_i) \rightarrow \max_{\beta'_0, \boldsymbol{\beta}}$$

where  $c_i$  is the class of  $x_i$ .

# Multiclass logistic regression

- Assumption:

$$\ln \left( \frac{\gamma_s p(\omega_s | x)}{\gamma_C p(\omega_C | x)} \right) = \beta_{s0} + \beta_s^T \mathbf{x}, \quad s = 1, 2, \dots, C-1$$

- Posterior class probabilities:

$$p(\omega_s | x) = \frac{\exp(\beta'_{s0} + \beta_s^T \mathbf{x})}{1 + \sum_{s=1}^{C-1} \exp(\beta'_{s0} + \beta_s^T \mathbf{x})}, \quad s = 1, 2, \dots, C-1$$

$$p(\omega_C | x) = \frac{1}{1 + \sum_{s=1}^{C-1} \exp(\beta'_{s0} + \beta_s^T \mathbf{x})}$$

$$\beta'_{s0} = \beta_{s0} - \ln(\gamma_s / \gamma_C)$$



## Multiclass logistic regression

- Decision rule (following Bayes minimum risk principle): assign  $x$  to class  $c = \arg \max_c \beta_{c0} + \beta_c^T x$  if  $\beta_{c0} + \beta_c^T x > 0$  otherwise assign  $x$  to class  $C$ .
- Estimate with ML:

$$\prod_{i=1}^n p(c_i | x_i) \rightarrow \max_{\beta'_0, \beta}$$

where  $c_i$  is the class of  $x_i$ .

- Please pay attention to the difference between  $\beta_0$  and  $\beta'_0$ .

# Logistic regression - loss function


For two class situation  $p(y|x) = \sigma(\langle w, x \rangle y)$  for  $\sigma = \frac{1}{1+e^{-z}}$ ,  
 $w = [\beta'_0, \beta]$ ,  $x = [1, x_1, x_2, \dots, x_D]$ .

Estimation with ML:

$$\prod_{i=1}^n \sigma(\langle w, x_i \rangle y_i) \rightarrow \max_w$$

which is equivalent to

$$\sum_i^n \ln(1 + e^{-\langle w, x_i \rangle y_i}) \rightarrow \min_w$$



It follows that logistic regression is linear discriminant estimated with loss function  $\mathcal{L}(M) = \ln(1 + e^{-M})$ .

## SGD realization of logistic regression


Substituting  $\mathcal{L}(M) = \ln(1 + e^{-M})$  into update rule, we obtain that for each sample  $(x_i, y_i)$  weights should be adapted according to

$$w \leftarrow w + \eta \sigma(-M_i) x_i y_i$$

Perceptron of Rosenblatt update rule:

$$w \leftarrow w + \eta \mathbb{I}[M_i < 0] x_i y_i$$

- Logistic rule update is the smoothed variant of perceptron's update.
- The more severe the error (according to margin) - the more weights are adapted.



## Logistic regression - assumptions

In logistic regression it was assumed that (for equal misclassification costs):

$$\ln \left( \frac{p(\omega_1|x)}{1 - p(\omega_1|x)} \right) = \beta_0 + \beta^T \mathbf{x}$$

which is equivalent to

$$p(\omega_1|x) = \frac{\exp(\beta_0 + \beta^T \mathbf{x})}{1 + \exp(\beta_0 + \beta^T \mathbf{x})}$$

Decision rule (following Bayes minimum risk principle):

$$\mathbf{x} = \begin{cases} \omega_1, & \beta_0 + \beta^T \mathbf{x} > 0 \\ \omega_2, & \beta_0 + \beta^T \mathbf{x} < 0 \end{cases}$$

What assumption allowed to obtain probabilities?

## Logistic regression - assumptions

In logistic regression it was assumed that (for equal misclassification costs):

$$F(p(\omega_1|x)) = \beta_0 + \boldsymbol{\beta}^T \mathbf{x}$$

Any  $F(z)$  satisfying:

- $F(z)$  is increasing
- $\text{Dom}[F] = (0, 1)$
- $\text{Im}[F] = \mathbb{R}$
- $F(1/2) = 0$

leads to the same decision rule:

$$\mathbf{x} = \begin{cases} \omega_1, & \beta_0 + \boldsymbol{\beta}^T \mathbf{x} > 0 \\ \omega_2, & \beta_0 + \boldsymbol{\beta}^T \mathbf{x} < 0 \end{cases}$$

## Logistic regression - assumptions

This is equivalent to

$$p(\omega_1|x) = G(\beta_0 + \beta^T \mathbf{x})$$

for any  $G = F^{-1}$ , satisfying:


- $G(z)$  is increasing
- $\text{Dom}[G] = \mathbb{R}$
- $\text{Im}[G] = (0, 1)$
- $G(0) = 1/2$

leads to the same decision rule:

$$\mathbf{x} = \begin{cases} \omega_1, & \beta_0 + \beta^T \mathbf{x} > 0 \\ \omega_2, & \beta_0 + \beta^T \mathbf{x} < 0 \end{cases}$$

# Probit/Logit

- $G(z)$  may be distribution function of any continuous symmetrical random variable, taking values on  $\mathbb{R}$ .
- Examples:
  - $G(z) = \frac{e^z}{1+e^z}$  - logit (leads to logistic regression)
  - $G(z) = \Phi(z)$  - normal c.d.f., probit.



# Analysis of logistic regression

## Advantages

- Implements margin offset strategy (using smoothed weights adaptation in SGD)
- Gives estimates of class probabilities



# Analysis of logistic regression

## Advantages

- Implements margin offset strategy (using smoothed weights adaptation in SGD)
- Gives estimates of class probabilities

## Disadvantages


- Quality deteriorates if model assumptions are violated
- Disadvantages inherited from SGD:
  - need to normalize features
  - filter outliers
  - regularization for multiple or correlated features

# Table of Contents



- 1 Stochastic gradient descent
- 2 Regularization
- 3 Logistic regression
- 4 Support vector machines**
- 5 Kernel support vector machines

## Reminder



For linear discriminant function  $g(x) = w^T x + w_0$ :



# Support vector machines



# Support vector machines



## Main idea

Select hyperplane maximizing the margin - the sum of distances from nearest  $\omega_1$  object to hyperplane and from nearest  $\omega_2$  object to hyperplane.

## Support vector machines

Objects  $x_i$  for  $i = 1, 2, \dots, n$  lie at distance  $b/|w|$  from discriminant hyperplane if

$$\begin{cases} x_i^T w + w_0 \geq b, & y_i = +1 \\ x_i^T w + w_0 \leq b & y_i = -1 \end{cases} \quad i = 1, 2, \dots, n.$$

This can be rewritten as

$$y_i(x_i^T w + w_0) \geq b, \quad i = 1, 2, \dots, n.$$

The margin is equal to  $2b/|w|$ . Since  $w$ ,  $w_0$  and  $b$  are defined up to multiplication constant, we can set  $b = 1$ .

## Problem statement

Problem statement:

$$\begin{cases} w^T w \rightarrow \min_{w, w_0} \\ y_i(x_i^T w + w_0) \geq 1, \quad i = 1, 2, \dots, n. \end{cases}$$

According to Karush-Kuhn-Takker theorem, solution satisfies the following problem:


$$L_P = \frac{1}{2} w^T w - \sum_{i=1}^n \alpha_i (y_i (w^T x_i + w_0) - 1) \rightarrow \min_{w, w_0} \max_{\alpha} \quad \alpha_i \geq 0, \quad i = 1, 2, \dots$$

with the constraints:

$$\begin{cases} \alpha_i \geq 0, \\ y_i(x_i^T w + w_0) - 1 \geq 0, \\ \alpha_i (y_i(x_i^T w + w_0) - 1) = 0. \end{cases}$$

# Support vectors

Condition  $\alpha_i(y_i(x_i^T w + w_0) - 1) = 0$  is satisfied when either  $\alpha_i = 0$  or  $y_i(x_i^T w + w_0) - 1 = 0$ . Second case describes support vectors, which lie at distance  $1/|w|$  to separating hyperplane and which affect the weights. Other vectors don't affect the solution.





## Dual problem

$$\frac{\partial L}{\partial w_0} = 0 : \sum_{i=1}^n \alpha_i y_i = 0$$

$$\frac{\partial L}{\partial w} = 0 : w = \sum_{i=1}^n \alpha_i y_i x_i$$

Substituting into Lagrangian  $L_P$ , we get:

$$L_D = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j \rightarrow \max_{\alpha}$$

$\alpha_i$  can be found from the dual optimization problem:

$$\begin{cases} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j \rightarrow \max_{\alpha} \\ \alpha_i \geq 0, i = 1, 2, \dots, n; \sum_{i=1}^n \alpha_i y_i = 0 \end{cases}$$

## Solution

Denote  $\mathcal{SV}$  - the set of indexes of support vectors.

Optimal  $\alpha_i$  determine weights directly:

$$w = \sum_{i \in \mathcal{SV}} \alpha_i y_i x_i$$

$w_0$  can be found from any edge equality for support vectors:

$$y_i(x_i^T w + w_0) = 1, i \in \mathcal{SV}$$

Solution from summation over  $n_{\mathcal{SV}}$  equation provides a more robust estimate of  $w_0$ :


$$n_{\mathcal{SV}} w_0 + \sum_{i \in \mathcal{SV}} x_i^T w = \sum_{i \in \mathcal{SV}} y_i$$

## Linearly non-separable case

No separating hyperplane exists. Errors are permitted by including slack variables  $\xi_i$ :

$$\begin{cases} \frac{1}{2} w^T w + C \sum_{i=1}^n \xi_i \rightarrow \min_{w, \xi} \\ y_i(w^T x_i + w_0) \geq 1 - \xi_i, i = 1, 2, \dots, n \\ \xi_i \geq 0, i = 1, 2, \dots, n \end{cases}$$

- Parameter  $C$  is the cost for misclassification and controls the bias-variance trade-off.
- It is chosen on validation set.
- Other penalties are possible, e.g.  $C \sum_i \xi_i^2$ .



## Linearly non-separable case

According to Karush-Kuhn-Takker theorem, the solution satisfies:

$$L_P = \frac{1}{2} w^T w + C \sum_i \xi_i - \sum_{i=1}^n \alpha_i (y_i (w^T x_i + w_0) - 1 + \xi_i) - \sum_{i=1}^n r_i \xi_i$$

$$L_P \rightarrow \min_{w, w_0, \xi} \max_{\alpha, r}$$

under constraints:

$$\begin{cases} \xi_i \geq 0, \alpha_i \geq 0, r_i \geq 0 \\ y_i (x_i^T w + w_0) \geq 1 - \xi_i, \\ \alpha_i (y_i (w^T x_i + w_0) - 1 + \xi_i) = 0 \\ r_i \xi_i = 0 \end{cases}$$

$$\frac{\partial L_P}{\partial \xi_i} = 0 : C - \alpha_i - r_i = 0 \quad \Rightarrow \quad \alpha_i \in [0, C].$$

# Classification of training objects

- **Non-informative objects:**

- have  $\alpha_i = 0$  ( $\Leftrightarrow r_i = C \Leftrightarrow \xi_i = 0 \Leftrightarrow y_i(w^T x_i + w_0) \geq 1$ )

- **Support vectors:**

- have  $\alpha_i > 0$  ( $\Leftrightarrow y_i(w^T x_i + w_0) = 1 - \xi_i$ )

- **boundary support vectors:**

- have  $\xi_i = 0$  ( $\Leftrightarrow r_i > 0 \Leftrightarrow \alpha_i \in (0, C) \Leftrightarrow y(w^T x_i + w_0) = 1$ )  
then support vector lies at  $1/|w|$  distance to separating hyperplane and is called boundary support vector.

- **violating support vectors:**

- have  $\xi_i > 0$  ( $\Leftrightarrow r_i = 0 \Leftrightarrow \alpha_i = C$ ), so lies closer than  $1/|w|$  to separating hyperplane.
- If  $\xi_i \in (0, 1)$  then violating support vector is correctly classified.
- If  $\xi_i > 1$  then violating support vector is misclassified.

## Linearly non-separable case - dual problem

$$\frac{\partial L_P}{\partial w_0} = 0 : \sum_{i=1}^n \alpha_i y_i = 0$$

$$\frac{\partial L_P}{\partial w} = 0 : w = \sum_{i=1}^n \alpha_i y_i x_i$$

$$\frac{\partial L_P}{\partial \xi_i} = 0 : C - \alpha_i - r_i = 0$$

Substituting these constraints into  $L_P$ , we obtain the dual problem:

$$\begin{cases} L_D = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j \rightarrow \max_{\alpha} \\ \sum_{i=1}^n \alpha_i y_i = 0 \\ 0 \leq \alpha_i \leq C \end{cases}$$

## Solution

Denote  $\mathcal{SV}$  - the set of indexes of support vectors with  $\alpha_i > 0$   
 $(\Leftrightarrow y(w^T x_i + w_0) = 1 - \xi_i)$  and  $\widetilde{\mathcal{SV}}$  - the set of indexes of support  
 vectors with  $\alpha_i \in (0, C)$   $(\Leftrightarrow \xi_i = 0, y(w^T x_i + w_0) = 1)$   
 Optimal  $\alpha_i$  determine weights directly:

$$w = \sum_{i \in \mathcal{SV}} \alpha_i y_i x_i$$

$w_0$  can be found from any edge equality for support vectors, having  
 $\xi_i = 0$ :

$$y_i(x_i^T w + w_0) = 1, i \in \widetilde{\mathcal{SV}}$$

Solution from summation of equations for each  $i \in \widetilde{\mathcal{SV}}$  provides a  
 more robust estimate of  $w_0$ :

$$n_{\widetilde{\mathcal{SV}}} w_0 + \sum_{i \in \widetilde{\mathcal{SV}}} x_i^T w = \sum_{i \in \widetilde{\mathcal{SV}}} y_i$$


## Another view on SVM

Optimization problem:

$$\begin{cases} \frac{1}{2} w^T w + C \sum_{i=1}^n \xi_i \rightarrow \min_{w, \xi} \\ y_i(w^T x_i + w_0) = M_i(w, w_0) \geq 1 - \xi_i, \\ \xi_i \geq 0, i = 1, 2, \dots, n \end{cases}$$

can be rewritten as

$$\frac{1}{2C} |w|^2 + \sum_{i=1}^n [1 - M_i(w, w_0)]_+ \rightarrow \min_{w, \xi}$$



Thus SVM is linear discriminant function with cost approximated with  $\mathcal{L}(M) = [1 - M]_+$  and  $L_2$  regularization.



## Probabilistic interpretation

SVM optimization task may be obtained if

$$p(x_i, y_i | w, w_0) \sim z_1 e^{-[1 - M_i(w, w_0)]_+}$$

with Gaussian prior probability

$$p(w | C) = z_2 e^{-|w|^2 / (2C)}$$

# Properties

Solution:

$$y = \text{sign} \left\{ \sum_{i \in \mathcal{SV}} \alpha_i y_i < x_i, x > + w_0 \right\}$$

Sparsity of SVM: solution depends only on support vectors:

- more affected by outliers
- possible filtering scheme (like editing):
  - solve
  - remove lowest margin objects
  - solve on refined sample
- Relevant vectors machine: filter support vectors by regularization of  $\alpha \sim \frac{1}{(2\pi)^{n/2} \sqrt{C_1 C_2 \dots C_n}} e^{(-\sum_{i=1}^n \frac{\alpha_i^2}{2C_i})}$
- if only a small fraction of objects are incorrectly classified, they may be removed from the training sample and it becomes separable  $\Rightarrow$  no need to select C.

## Multiclass classification

$C$  classes  $\omega_1, \omega_2, \dots, \omega_C$ .

- One-against-all:
  - build  $C$  binary classifiers, classifying class  $\omega_i$  against other classes
  - select the class with highest margin
- One-against-one:
  - build  $C(C-1)/2$  classifiers, classifying class  $\omega_i$  against  $\omega_j$ .
  - select the class having maximum votes
- Multiclass variant of initial algorithm

# Multiclass SVM

$C$  discriminant functions are built simultaneously:

$$g_k(x) = (w^k)^T x + w_0^k$$

Linearly separable case:

$$\begin{cases} \sum_{k=1}^C (w^k)^T w^k \rightarrow \min_w \\ (w^{c(i)})^T x + w_0^{c(i)} - (w^k)^T x - w_0^k \geq 1 \quad \forall k \neq c(i), i = 1, 2, \dots, n \end{cases}$$

Linearly non-separable case:

$$\begin{cases} \sum_{k=1}^C (w^k)^T w^k + C \sum_{i=1}^n \xi_i \rightarrow \min_w \\ (w^{c(i)})^T x + w_0^{c(i)} - (w^k)^T x - w_0^k \geq 1 - \xi_i \quad \forall k \neq c(i), i = 1, 2, \dots, n \\ \xi_i \geq 0 \end{cases}$$

# Table of Contents

- 1 Stochastic gradient descent
- 2 Regularization
- 3 Logistic regression
- 4 Support vector machines
- 5 Kernel support vector machines**

## Linear SVM reminder

- Solution for weights:

$$w = \sum_{i \in \mathcal{SV}} \alpha_i y_i x_i$$

Discriminant function

$$g(x) = \sum_{i \in \mathcal{SV}} \alpha_i y_i \langle x_i, x \rangle + w_0$$

$$w_0 = \frac{1}{n_{\widetilde{\mathcal{SV}}}} \left( \sum_{i \in \widetilde{\mathcal{SV}}} y_i - \sum_{i \in \widetilde{\mathcal{SV}}} \sum_{j \in \mathcal{SV}} \alpha_j y_j \langle x_i, x_j \rangle \right)$$

## Kernel SVM

- $x$  is replaced with  $\phi(x)$
- $[x] \rightarrow [x, x^2, x^3]$

### Kernel

Function  $K(x, y) : X \times X \rightarrow \mathbb{R}$  is a kernel function if it may be represented as  $K(x, y) = \langle \psi(x), \psi(y) \rangle$  for some mapping  $\psi : X \rightarrow H$ , with scalar product defined on  $H$ .

- $\langle x, y \rangle$  is replaced by  $\langle \phi(x), \phi(y) \rangle = K(x, y)$

## Kernel SVM

Discriminant function

$$g(x) = \sum_{i \in \mathcal{SV}} \alpha_i y_i K(x_i, x) + w_0$$

$$w_0 = \frac{1}{n_{\widetilde{\mathcal{SV}}}} \left( \sum_{i \in \widetilde{\mathcal{SV}}} y_i - \sum_{i \in \widetilde{\mathcal{SV}}} \sum_{j \in \mathcal{SV}} \alpha_j y_j K(x_i, x_j) \right)$$



## Kernel properties

**Theorem (Mercer):** Function  $K(x, y)$  is a kernel is and only if

- it is symmetric:  $K(x, y) = K(y, x)$
- it is non-negative definite: for every function  $g : X \rightarrow \mathbb{R}$

$$\int_X \int_X K(x, x') g(x) g(x') dx dx' \geq 0$$

- Example:  $K(x, y) = (1 + x^T y)^2 = (1 + x_1 y_1 + x_2 y_2)^2 = 1 + 2x_1 y_1 + 2x_1 y_2 + 2x_1 x_2 y_1 y_2 + x_1^2 y_1^2 + x_2^2 y_2^2 = \phi^T(x) \phi(x)$
- $\phi(x) = (1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1 x_2, x_1^2, x_2^2)$

## Kernel properties

Kernels can be constructed manually:


- Scalar product  $\langle x, x' \rangle$  is a kernel
- Constant  $K(x, x') \equiv 1$  is a kernel
- Product of kernels  $K(x, x') = K_1(x, x')K_2(x, x')$  is a kernel.
- For every function  $\psi : X \rightarrow \mathbb{R}$  the product  $K(x, x') = \psi(x)\psi(x')$  is a kernel
- Linear combination of kernels  $K(x, x') = \alpha_1 K_1(x, x') + \alpha_2 K_2(x, x')$  with positive coefficients is a kernel
- Composition of function  $\varphi : X \rightarrow X$  and kernel  $K_0$  is a kernel:  $K(x, x') = K_0(\varphi(x), \varphi(x'))$
- etc.

Useful collection of datasets matching datasets and research papers:  
<https://archive.ics.uci.edu/ml/datasets.html>





## Commonly used kernels


Let  $x$  and  $y$  be two objects.

Kernel	Mathematical form
linear	$\langle x, y \rangle$
polynomial	$(\gamma \langle x, y \rangle + r)^d$
RBF	$\exp(-\gamma  x - y ^2)$
sigmoid	$\tanh(\gamma \langle x, y \rangle + r)$





## Kernel results




Linear kernel - variable  $C$ 

## Linear kernel - variable C




Linear kernel - variable  $C$ 


Linear kernel - variable  $C$ 





## RBF kernel - variable $\gamma$




## RBF kernel - variable $\gamma$




RBF kernel - variable  $\gamma$ 

RBF kernel - variable  $\gamma$ 


## RBF kernel - variable C




## RBF kernel - variable C



## RBF kernel - variable C




# Polynomial kernel - variable $d$







## Polynomial kernel - variable $d$




# Polynomial kernel - variable $d$




# Polynomial kernel - variable $d$




## Polynomial kernel - variable $C$




## Polynomial kernel - variable C




## Polynomial kernel - variable C




Sigmoid kernel - variable  $\gamma$ 

## Sigmoid kernel - variable $\gamma$





## Sigmoid kernel - variable $\gamma$



Sigmoid kernel - variable  $C$ 